

PROJECT SPECIFICATION

Cloud Capstone Project

(Option 1): CI/CD, Github & Code Quality

CRITERIA	MEETS SPECIFICATIONS
The project demonstrates an understanding of CI and Github.	All project code is stored in a GitHub repository and a link to the repository has been provided for reviewers. The student uses a CI tool to build the application
The project has a proper documentation.	The README file includes introduction how to setup and deploy the project. It explains the main building blocks and has comments in the important files.
The project use continuous deployments (CD)	A CD tool is in place to deploy new version of the app automatically to production. The way is described and easy to follow.

(Option 1): Container

CRITERIA	MEETS SPECIFICATIONS
The app is containerized	There is a Dockerfile in repo and the docker image can be build
The project have public docker images	On DockerHub images for the application are available
The applications runs in a container without errors	Starting the app as a container on a local system

(Option 1): Deployment

CRITERIA	MEETS SPECIFICATIONS
The application runs on a cluster in the cloud	The project can be deployed to a kubernetes cluster

CRITERIA	MEETS SPECIFICATIONS
The app can be upgraded via rolling-update	The students can deploy a new version of the application without downtime
A/B deployment of the application	Two versions of the same app can run at the same and service traffic
Monitoring	The application is monitored by Amazon CloudWatch

(Option 2): Functionality

CRITERIA	MEETS SPECIFICATIONS
The application allows users to create, update, delete items	A user of the web application can use the interface to create, delete and complete an item.
The application allows users to upload a file.	A user of the web interface can click on a "pencil" button, then select and upload a file. A file should appear in the list of items on the home page.
The application only displays items for a logged in user.	If you log out from a current user and log in as a different user, the application should not show items created by the first account.
Authentication is implemented and does not allow unauthenticated access.	A user needs to authenticate in order to use an application.

(Option 2):Codebase

CRITERIA	MEETS SPECIFICATIONS
The code is split into multiple layers separating business logic from I/O related code.	Code of Lambda functions is split into multiple files/classes. The business logic of an application is separated from code for database access, file storage, and code related to AWS Lambda.

CRITERIA	MEETS SPECIFICATIONS
Code is implemented using async/await and Promises without using callbacks.	To get results of asynchronous operations, a student is using async/await constructs instead of passing callbacks.

(Option 2):Best practices

CRITERIA	MEETS SPECIFICATIONS
All resources in the application are defined in the "serverless.yml" file	All resources needed by an application are defined in the "serverless.yml". A developer does not need to create them manually using AWS console.
Each function has its own set of permissions.	Instead of defining all permissions under provider/iamRoleStatements , permissions are defined per function in the functions section of the "serverless.yml".
Application has sufficient monitoring.	<p>Application has at least some of the following:</p> <ul style="list-style-type: none"> • Distributed tracing is enabled • It has a sufficient amount of log statements • It generates application level metrics
HTTP requests are validated	Incoming HTTP requests are validated either in Lambda handlers or using request validation in API Gateway. The latter can be done either using the serverless-request-validator-plugin or by providing request schemas in function definitions.

(Option 2):Architecture

CRITERIA	MEETS SPECIFICATIONS
Data is stored in a table with a composite key.	<p>1:M (1 to many) relationship between users and items is modeled using a DynamoDB table that has a composite key with both partition and sort keys. Should be defined similar to this:</p> <pre> KeySchema: - AttributeName: partitionKey KeyType: HASH - AttributeName: sortKey KeyType: RANGE </pre>

CRITERIA	MEETS SPECIFICATIONS
Scan operation is not used to read data from a database.	Items are fetched using the "query()" method and not "scan()" method (which is less efficient on large datasets)
